School of Science and Engineering

Title

Capstone Design Interim Report

March 2025

Student Name

Supervised by:

Dr. Tajj-Eddine Rachidi

© SCHOOL OF SCIENCE & ENGINEERING – AL AKHAWAYN UNIVERSITY

Student conduct and copyright

The supervisor and the student (the Capstoner) agree that:

1. Permission has been obtained for any third party content (eg Data, illustrations, photographs, charts or maps).

2. The results described in this report have not previously been published

Student's name and signature: **Adnane Ahroum** 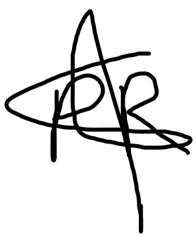

Supervisor's name and signature: **Dr.Tajj-Eddine Rachidi**

Deep Learning-Driven MRI Analysis for Precise Brain Tumor Classification, Detection, and Segmentation.

Capstone Report

Student Statement:

This capstone project addresses critical challenges in brain tumor diagnostics within Morocco's resource-constrained healthcare system, which is hindered by outdated infrastructure, high costs, and limited access to specialists. By integrating AI-driven medical imaging technologies, the project employs three key models: U-Net for tumor segmentation, ResNet-50 for subtype classification, and Faster R-CNN for detection. Utilizing the "Decathlon dataset" and simulating low-resolution MRI data, the framework is optimized for computational efficiency, ensuring feasibility on low-end hardware. Developed with Python, MONAI, and TensorFlow, this solution will aim to significantly reduce diagnostic costs, improve accuracy, and enhance healthcare accessibility in Morocco.

____

Student's name
Adnane Ahroum

Approved by the Supervisor

____

Dr. Tajj-Eddine Rachidi

ACKNOWLEDGEMENTS

I would like to express my sincere gratitude to Dr. Tajjedine Rachidi, whose Introduction to AI course first sparked my interest in the field. His teaching and insights have been a key source of motivation, inspiring me to explore the intersection of AI and healthcare for my master's studies and future research endeavors. I am also deeply thankful to him for his invaluable guidance and expertise as my supervisor, providing unwavering support throughout this project. Finally, I extend my heartfelt appreciation to my parents, whose encouragement continuously pushes me toward greater initiatives and personal growth.

CONTENT

LIST OF ACRONYMS & KEYWORDS

MRI: Magnetic Resonance Imaging.

CNN: Convolutional Neural Network**.**

ReLu: Rectified Linear Unit**.**

IEEE: Institute of Electrical and Electronics Engineers**.**

ISO: International Organization for Standardization**.**

CAD : Computer-Aided Diagnosis.

MONAI: Medical Open Network for AI.

mAP: mean Average Precision.

MoF: Ministry of Health.

DICOM: Digital Imaging and Communications in Medicine.

PACS: Picture Archiving and Communication Systems

LIST OF FIGURES

LIST OF TABLES

ABSTRACT (ENGLISH):

Brain tumors pose significant health challenges, particularly in resource-limited settings like Morocco. This project aims to develop an affordable, AI-driven system to enhance brain tumor diagnosis. By integrating advanced deep learning models and utilizing accessible MRI datasets, we seek to improve diagnostic accuracy and accessibility in Moroccan healthcare. This initiative has the potential to reduce reliance on limited radiology resources and positively impact patient outcomes through earlier and more reliable diagnoses.

Keywords: Brain Tumors, AI-Driven Diagnosis, Deep Learning Models, Moroccan Healthcare

RESUME (FRENCH) :

Les tumeurs cérébrales constituent des défis sanitaires majeurs, notamment dans des contextes à ressources limitées comme le Maroc. Ce projet vise à développer un système abordable, basé sur l'intelligence artificielle, pour améliorer le diagnostic des tumeurs cérébrales. En intégrant des modèles avancés d'apprentissage profond et en utilisant des ensembles de données IRM accessibles, nous cherchons à accroître la précision diagnostique et l'accessibilité des soins de santé au Maroc. Cette initiative pourrait réduire la dépendance aux ressources limitées en radiologie et améliorer les résultats pour les patients grâce à des diagnostics plus précoces et plus fiables.

Mots clés : Tumeurs cérébrales, Systèmes basé sur IA, Diagnostic médical au maroc, Modèles d’apprentissage profond

1 Introduction:

Brain tumors are among the most dangerous forms of cancer, where early and accurate diagnosis is vital for effective treatment. In many low- and middle-income countries, including Morocco, timely diagnosis is often hampered by limited medical infrastructure and specialist availability. The availability of advanced diagnostic tools such as MRI is far below international standards – for example, Morocco has only around 0.36 MRI units per million people, trailing behind neighboring countries (Ogbole et al., 2018). This scarcity means patients in rural or underserved areas may face long wait times or travel great distances for neuroimaging. Additionally, there is a shortage of expert neuroradiologists, and existing diagnostic equipment is frequently outdated and costly, making modern brain tumor evaluation inaccessible to many. Another challenge in Morocco's healthcare system is the inaccuracy or inconsistency of current diagnostic methods. Physicians must often rely on subjective image interpretations, which can lead to missed early-stage tumors or misclassification of tumor types. Confirmatory diagnosis via biopsy is invasive and not always immediately available. Furthermore, while AI-based computer-aided diagnosis (CAD) tools for brain tumors have shown promise globally, local adoption has been slow. Medical professionals have been hesitant to incorporate AI into the diagnostic process due to earlier systems' reliability issues. In under-resourced settings, there is an urgent need for affordable, reliable AI solutions that can assist clinicians by highlighting tumor regions and suggesting probable classifications, thus acting as a "second pair of eyes." This could democratize access to quality diagnostics, alleviating inequalities in care.

2. Problem Statement:

Current brain tumor diagnostic methods in Morocco face several interrelated issues of cost, accessibility, and accuracy. Firstly, advanced imaging like MRI is expensive to obtain and maintain. The high cost of MRI scanners and related infrastructure means they are concentrated in a few urban centers, limiting accessibility for most of the population. Patients from remote or low-income regions may not get scanned until the disease is advanced. This gap is evidenced by Morocco's limited access to advanced diagnostics. Secondly, the scarcity of equipment and specialists leads to delays in diagnosis, adversely affecting outcomes for brain tumor patients. Early symptoms might be overlooked or attributed to other causes due to lack of immediate imaging, and when imaging is done, backlogs in radiology departments can further slow the process.

Additionally, even when imaging is available, the accuracy of diagnosis can be constrained by human and technical factors. Radiological expertise varies, and subtle tumors or distinctions between tumor types (e.g., differentiating a low-grade glioma from a benign meningioma) may be missed. Pathological confirmation through biopsy is the gold standard but involves surgical risks and time. There is also limited use of sophisticated image analysis (such as volumetric measurements or computer-assisted detection) in the current system, due in part to outdated tools and lack of AI integration. The result is that diagnostic errors or uncertainties can occur, and treatment planning is often initiated with incomplete information.

Moreover, the existing approach is costly for both patients and the healthcare system. An MRI scan and subsequent specialist consultation can be prohibitively expensive for many Moroccan patients, leading to financial strain or avoidance of care altogether. Hospitals bear high operational costs for running scanners and importing imaging supplies, which is unsustainable in a resource-limited context. In summary, this project directly addresses these issues: it proposes an AI-driven solution to lower costs (by using open-source tools and potentially enabling diagnosis on widely available hardware), increase accessibility (through automated analysis that can reach hospitals lacking specialists), and improve diagnostic accuracy (via proven deep learning models that can highlight tumors and predict their type). By clearly defining these problems, we establish the need for an innovative approach that leverages AI to augment the current diagnostic pathway.

3. Project Specifications:

Functional Overview: The project will deliver an AI-based pipeline capable of taking a patient's brain MRI scan and automatically segmenting tumor tissue, classifying the tumor type, and detecting its location within the brain. Ultimately, this could be integrated into a user-friendly application for radiologists and clinicians in Morocco.

- **AI Models**: Three complementary deep learning models are being developed:
    - **Segmentation Model (U-Net)**: Using the MONAI framework, a 3D U-Net architecture will segment brain tumors from MRI scans, delineating the tumor boundaries within each slice. The model is trained to maximize the Dice similarity coefficient for accurate tumor masks.
    - **Classification Model (ResNet-50)** : A CNN based on ResNet-50 (adapted for 3D image inputs) will classify the segmented tumor into categories such as Meningioma, Glioma, or Pituitary tumor. Transfer learning is employed by initializing with ImageNet weights to improve performance on the relatively small medical dataset. The target is to achieve high accuracy in line with literature (over 90% classification accuracy).
    - **Detection Model (Faster R-CNN)**: A region-proposal based object detection model (implemented in PyTorch) will scan entire MRI volumes to localize tumors, outputting bounding boxes around tumor regions. This provides an independent check to ensure even small or subtle tumors are flagged, complementing the

segmentation output. The detector's parameters (such as anchor sizes) are tuned for brain anatomy to reduce false detections.

- **Tools & Libraries**: The development leverages a suite of open-source tools well-suited for medical imaging. MONAI (Medical Open Network for AI) is used for building and training the U-Net segmentation model, given its specialized utilities for 3D medical image processing. 3D Slicer (an open-source medical image analysis platform) is utilized for data preprocessing. For the classification and detection models, deep learning frameworks such as PyTorch and TensorFlow are employed: PyTorch (which underpins MONAI) is mainly used for segmentation and detection, while TensorFlow/Keras is an option for implementing the ResNet-50 classification model (ensuring experience with both major frameworks). The training and experimentation environment is Python-based, with Jupyter notebooks on platforms like Google Colab to take advantage of GPU acceleration.

4. STEEPLE Analysis:

- **Social:** The project has a clear positive social impact by improving access to early brain tumor diagnosis in under-resourced Moroccan hospitals. Currently, patients in rural or public healthcare settings often cannot get timely expert diagnoses. An AI tool that assists doctors can bridge this gap, leading to earlier detection and treatment. By reducing the need for specialist interpretation on-site, patients in remote regions could have their MRI scans analyzed with expert-level insight. Additionally, empowering doctors with AI may increase patient trust in diagnoses; families and communities' benefit when serious conditions are identified and treated sooner. A potential social challenge is the **acceptance of AI in healthcare**: both patients and some medical staff may be initially hesitant to trust AI recommendations. The project addresses this by positioning AI as a decision-support tool for doctors (not a replacement), emphasizing improved outcomes.

- **Technological:** Technologically, this project leverages state-of-the-art deep learning techniques for medical imaging, contributing to the modernization of diagnostics in Morocco. It reduces dependency on expensive, fixed diagnostic tools by enabling software-driven analysis on standard computers. The technology is built on open-source frameworks, which ensure adaptability and community support. A consideration here is the local infrastructure: hospitals must have at least basic computing hardware (preferably cloud access) to run the models. Another factor is internet connectivity if the tool were cloud-based; however, the overall implementation can be kept local to avoid connectivity issues and protect data. The risk of obsolescence is mitigated by the project's use of proven architecture, which can be updated as needed. Moreover, by using **transfer learning** and pre-trained models, the project adapts advanced technology to relatively small datasets effectively.
- **Economic:** Economically, the AI solution is poised to reduce healthcare costs for both providers and patients. By improving diagnostic efficiency, unnecessary procedures or delays can be minimized, thus cutting expenses. Early and accurate detection of tumors can also reduce the cost burden of treating advanced cancers. The system itself is designed to be low-cost: it uses open-source software and commodity hardware, making it feasible for public hospitals with limited budgets. This addresses the economic barrier where advanced diagnostics were once only available to those who could afford private care. Lowering costs for patients (e.g. by avoiding multiple specialist consultations or long trip for second opinions) aligns with the project's goal. In the long run, healthier patients contribute more to the economy, and the healthcare system saves resources by catching illnesses early. The project might require an initial investment in training staff and setting up the system, but this is marginal compared to the potential savings. There is also an opportunity for economic support from government or partnerships (as noted, the Health

Ministry or universities might collaborate), which could fund the scale-up if the pilot is successful.

- **Environmental:** The environmental impact of this project is minimal. Unlike large medical devices or new facility construction, an AI software solution has a very small environmental footprint. It primarily requires electricity for computing. By potentially reducing the need for patients to travel long distances for diagnosis (through tele-radiology use of the AI), the project can indirectly cut down on travel-related carbon emissions. The use of existing MRI machines more efficiently (analyzing digital images) has no additional environmental cost. One consideration is the energy consumption of GPUs during model training and deployment, but this is relatively minor and can be optimized. No hazardous waste or significant resource consumption is involved.
- **Political:** This project aligns with national healthcare improvement goals and thus holds neutral to positive political implications. There is growing political will in Morocco to modernize healthcare and integrate digital solutions. The AI-driven diagnostic tool could be seen as a step toward innovation and better public health services, supporting government initiatives. It does not directly interfere with any political or religious sensitivities as it is a medical tool. However, political support may be needed for wide adoption, such as endorsement by the MoF or inclusion in national healthcare plans. Successful implementation could bolster confidence in public hospitals, a politically favorable outcome. The project itself is apolitical; it serves patients regardless of background. We do note that sustained support (funding for maintenance, integration into hospital IT systems) might require advocacy and policy decisions.

- **Legal:** Implementing an AI solution in healthcare necessitates careful attention to legal requirements, especially concerning patient data. In Morocco, **Law 09-08** on personal data protection governs the handling of medical data. The project must ensure full compliance with this law, which means that MRI scans and any patient information used to train or run the models must be properly anonymized and secured. Data agreements with hospitals need to clarify usage rights and ensure no breach of confidentiality. Additionally, if the AI system is considered a medical device or diagnostic aid, there may be regulatory approval processes to navigate (though currently Morocco is developing frameworks for AI in medicine). We have also considered international standards like GDPR as good practice, even if not directly enforceable, to maintain high data privacy standards. Another legal aspect is liability: decisions made with AI assistance should still be validated by qualified professionals to avoid malpractice issues. Clear disclaimers and usage guidelines will be in place, aligning with global practices for AI decision support tools. In summary, the legal environment demands robust data security and privacy measures (which we adhere to), and appropriate regulatory navigation, but does not pose any insurmountable barrier.
- **Ethical:** Ethically, the project strives to uphold the highest standards, given that it deals with life-altering medical decisions. Patient consent and privacy are paramount; patients should ideally be informed if AI is being used in their diagnostic process, and their data should only be used for stated purposes. The project promotes beneficence: using AI to do good by catching tumors early and accurately. However, we must also address potential ethical concerns: for instance, the risk of algorithmic bias. If the training data is not sufficiently representative of Moroccan patients (e.g., mostly from foreign datasets), the AI might perform less well on local cases, raising equity concerns. We mitigate this by incorporating synthetic “Moroccan-like” data and plan for continuous model evaluation in the local context. Transparency is another ethical imperative – the model’s suggestions should be explainable to some degree, or at least outputs (segmented images, etc.) should

be viewable by doctors to validate. We also consider the resistance to AI adoption from an ethical standpoint: some healthcare professionals may feel threatened by AI or worry it diminishes their role. It's important to involve doctors in the development and deployment, addressing their concerns and demonstrating that the tool improves patient care (e.g., by showing time/cost savings and better outcomes). Lastly, obtaining ethical approval and oversight, especially if patient data is used for model training, is part of the plan in accordance with research ethics boards.

5. Engineering Standards:

- **DICOM (Digital Imaging and Communications in Medicine):** DICOM is the standard format for handling, storing, and transmitting medical images in hospitals. Our system adheres to DICOM protocols for input and output of imaging data, which means it can seamlessly integrate with existing hospital PACS (Picture Archiving and Communication Systems). By using DICOM-compliant data structures, we ensure interoperability: MRIs can be fed into AI without special conversion, and the segmented images or overlays produced can be saved in standard formats to be reviewed by clinicians on any DICOM viewer. This compliance is crucial for real-world deployment, as compatibility with hospital systems avoids any proprietary barriers.
- **IEEE Standards for AI Model Development:** In developing the machine learning models, we follow best practices akin to those suggested by IEEE guidelines (for instance, IEEE 2811 for reliability of AI, or principles from IEEE 829 on testing). We ensure that the training data is managed properly and the models are evaluated with standard metrics for performance (accuracy, Dice score, etc.). Floating-point arithmetic issues are kept in check (consistent precision) possibly referencing IEEE 754 standard for computations. Moreover, any future networked data sharing or federated learning

(as an extension to gather data from multiple hospitals) would consider standards like IEEE 2621 (medical device cybersecurity) and others relevant to safe AI in healthcare (Technologies, 2025).

6. Logic Model Framework:

1.1. Target Population:

The primary beneficiaries of this project are Moroccan patients who are at risk of brain tumors or undergoing diagnostic evaluation. By extension, their families also benefit from earlier and more accurate diagnoses leading to timely treatment. Another key group served is Moroccan healthcare providers, including radiologists, neurologists, and oncologists in public hospitals, who will use the AI tool as a diagnostic aid. Hospitals and clinics that currently lack advanced diagnostic tools stand to gain significantly; for example, regional hospitals without neuroradiology specialists could still offer high-quality MRI analysis through this system. In summary, the project serves patients (improving health outcomes), doctors (providing decision support), and healthcare institutions (enhancing service capabilities) across Morocco, particularly in underserved areas.

1.2. Underlying Assumptions:

A few important assumptions underlie the success of this project. We assume that sufficient data will be available to train and validate the AI models. This includes publicly available MRI datasets and any additional synthetic or local data we generate. We also assume that the computational resources (GPU access, etc.) are adequate to develop the model. Another assumption is that AI model performance on the curated dataset will translate effectively to real-world settings; essentially, we presume the gap between research data and hospital data can be bridged (with proper augmentation). Importantly, we operate under the belief that AI can compensate for limited radiologist expertise in practice, i.e., the models will be accurate enough to genuinely aid diagnosis. We also assume openness of stakeholders to adopt new technology, which is bolstered if we demonstrate that open-source tools can reduce costs by 60–80% compared to proprietary solutions. These assumptions are continually evaluated as the project progresses (for example, if data scarcity becomes an issue, we adapt by focusing more on transfer learning and augmentation).

1.3. Resources/Challenges:

Key resources for the project include access to public MRI datasets (like the Decathlon/BraTS dataset) and the rich documentation and support for tools such as MONAI and 3D Slicer. We have student and faculty expertise in deep learning to leverage, and computing resources via cloud GPU instances (Google Colab) to train models. Collaborations or consultations with medical professionals provide domain guidance. However, there are notable challenges: data scarcity is a concern: while public data exists, obtaining local MRI scans from Moroccan hospitals would greatly enhance the model, but requires navigating privacy and bureaucracy (initial plans to partner with a hospital had to be shelved due to time constraints). Hardware limitations are another challenge; not all deployment sites will have powerful GPUs, so we must optimize models for inference on modest hardware. We also face the usual project timeline challenge, ensuring that all three AI components can be developed and tested within the capstone schedule. Risk mitigation strategies are in place for each challenge (for example, using pre-trained models to counter data scarcity, and model quantization to help with hardware). Resources and challenges are two sides of the project reality that we manage proactively.

1.4. Activities:

1.4.1. Data Collection & Preprocessing: Acquire MRI scans from the Decathlon dataset, process them using 3D Slicer to standardize dimensions, normalize intensities, reduce noise, and generate synthetic variations to enhance the training set.

1.4.2. Model Development:

- Segmentation: Implement a U-Net model with MONAI, training it on preprocessed MRI volumes and corresponding tumor masks, adjusting parameters as needed.
- Classification: Fine-tune a ResNet-50 model on features from segmented tumor regions using adequate frameworks, optimizing through hyperparameter tuning and cross-validation.
- Detection: Set up a Faster R-CNN architecture in PyTorch to process entire MRI slices, output bounding boxes, and train it using scans with labeled tumor regions derived from segmentation masks.

1.4.3. Integration & User Interface: Develop a simple interface or script to sequentially run the three models on new MRI scans, compiling results to display segmented tumors with labels, confidence scores, or bounding outlines.

1.4.4. Testing & Evaluation: Assess model performance on a reserved test set and synthetic cases representative of Moroccan patients; seek qualitative feedback from radiologists to judge outputs.

1.4.5. Documentation & Dissemination: Document code and provide guidelines **for** model usage; prepare reports or presentations, potentially proposing findings at conferences.

1.5. Outputs of the Project:

The tangible deliverables from the project will include both data and models. One key output is an open-source segmented brain tumor dataset tailored for Morocco. Essentially, any synthetic MRI data we generate plus the segmentation labels can be packaged (within ethical/legal bounds) for future research use. The core outputs are the trained AI models themselves: (a) the segmentation model, capable of producing tumor masks from MRIs, (b) the classification model, that outputs tumor type given an ROI or segmented region, and (c) the detection model, that marks tumor locations on MRI slices. Along with the models, a prototype software pipeline that combines these components is expected from loading an MRI to printing out the diagnosis aid. and a final report describing the design and results. Another output is performance metrics and validation results: for instance, confusion matrices for the classifier, and segmented images. These artifacts showcase the project's efficacy. Finally, we aim to produce a conference proposal or academic poster as an output (as noted in our plan), summarizing the project for a broader scientific audience. Collectively, these outputs demonstrate the technical achievements and provide the foundation for real-world deployment.

1.6. Outcomes:

If the project's outputs are utilized, we anticipate significant short-term and long-term outcomes that benefit healthcare. In the short term, a successful proof-of-concept will show that a low-cost diagnosis aid for brain tumors is feasible. This could encourage further investment and interest in AI solutions within Moroccan healthcare. Intermediate outcomes might include pilot implementations in a partner hospital: for example, testing the tool on retrospective patient scans to measure improvement in diagnostic agreement or speed. As confidence in the tool grows, a long-term outcome would be integration into routine clinical practice in Moroccan clinics and hospitals. This means radiologists or general doctors use the AI system as part of their workflow for patients who get MRI scans. The ultimate long-term outcome is improved patient health: earlier tumor detection leads to earlier treatment (surgery, chemotherapy, etc.), which can improve survival rates and quality of life. We also foresee capacity building as an outcome: training local staff in AI usage and perhaps stimulating further local innovation in medical AI. On the systemic level, the project could influence health policy (outcome in form of greater support for digital health initiatives.

7. Literature Review:

Early adoption of AI for brain tumor diagnosis focused on using traditional machine learning and simple neural networks to classify tumor images. However, those CAD systems often faced accuracy issues, which limited their acceptance by medical professionals. For instance, rule-based or shallow learning systems could not consistently distinguish tumor types, leading to hesitancy in clinical uptake. This spurred researchers to develop more robust deep learning models. Amin et al. (2022) proposed a deep learning framework for brain tumor classification that significantly improved accuracy by leveraging data augmentation and transfer learning. In their approach, an enhanced CNN was trained on MRI data with techniques like rotation and flipping to increase data diversity, and pre-trained model weights (ResNet-50, VGG-16, etc.) were fine-tuned for the task. The result was a model with high predictive accuracy, outperforming baseline models, and the authors advocated its use in practical diagnosis settings. This underscores a trend in literature: applying transfer learning to compensate for limited medical datasets. Transfer learning has indeed shown remarkable success; one study reported that using a pre-trained CNN and fine-tuning it for brain MRI classification achieved about *94.8% accuracy*, a level approaching expert human performance. Another study's integrated CNN model reached an accuracy of 97.4% on an augmented brain tumor MRI dataset, highlighting the potential of deep networks when sufficiently trained. These results guide our project's classification component, where we use ResNet-50 with transfer learning, expecting a substantial accuracy boost on our targeted classes.

In terms of brain tumor segmentation, the research community has coalesced around the use of U-Net and its variants as the de facto standard. The Brain Tumor Segmentation (BraTS) challenges (2015–2021) have provided common benchmarks for comparing algorithms. Studies report that modern 3D U-Net architectures can achieve high Dice similarity scores (a measure of overlap between automated segmentation and ground truth). For example, a DenseNet-encoder U-Net evaluated on the BraTS 2018 dataset achieved Dice scores of 0.90 for whole tumor and around 0.79–0.85 for specific sub-regions (Stawiaski 11 & 11).

Such high accuracy in segmentation is crucial because delineating the tumor's extent is often the first step in diagnosis and treatment planning. The literature also explores combining segmentation with classification: some researchers first segment the tumor region and then use that isolated region to classify tumor type, which is precisely the approach we adopt. Cheng et al., and others have noted that focusing on the tumor region (via segmentation) can improve classification performance by filtering out irrelevant brain structures.

Another relevant area of literature is object detection in medical images, though it's slightly less common than pure segmentation for tumors. Techniques like Faster R-CNN have been successfully used in pathology or in detecting other anomalies (like polyps in colonoscopy videos), and a few works applied them to localize brain tumors on scans. These studies typically convert the problem into 2D object detection on each slice or 3D bounding boxes in a volume. While segmentation provides pixel-level detail, detection frameworks are faster and can highlight approximate locations. Our inclusion of a detection model is partly inspired by such works and the need for a quick screening tool: literature suggests that a well-trained detector can rapidly scan through volumes and flag suspicious areas for further analysis, acting as a triage mechanism.

Regarding the use of MONAI and 3D Slicer, our project leverages insights from both the academic and developer communities. MONAI is a relatively recent framework (built on PyTorch) explicitly designed for medical imaging tasks. Tutorials and studies using MONAI show that it greatly simplifies implementing complex models like 3D U-Net and provides domain-specific data augmentation (randomizable flips, intensity adjustments on MRI, etc.). For example, a MONAI tutorial on brain tumor segmentation demonstrates end-to-end training on the BraTS dataset with high performance, indicating that using such a specialized toolkit can speed up development without sacrificing model quality. We consider these best practices, such as using MONAI's efficient loaders for volumetric data and its implementations of Dice loss, as guiding lights for our segmentation model design.

3D Slicer, on the other hand, has been a staple in medical image visualization and even manual segmentation for decades (Kikinis & Pieper, 2011). Literature and user case reports show 3D Slicer being used as a platform to apply AI models in a user-friendly way. For instance, IEEE publications have discussed Slicer modules that integrate AI segmentation, allowing clinicians to manually refine AI-generated contours. This informs our vision that in the future, our models could be packaged as a 3D Slicer extension (there is even an existing extension called "SlicerDeepSeg" for general segmentation). From the literature perspective, using 3D Slicer for preprocessing is well-founded – studies indicate that consistent preprocessing (e.g., skull-stripping, normalization) significantly improves AI model results, and Slicer provides tools for these steps with a visual validation. We align with that knowledge by using Slicer to prepare our training data and generate synthetic images.

Finally, previous works specifically addressing the Moroccan or North African context are sparse, highlighting a gap that our project addresses. A review of healthcare literature shows discussions about healthcare disparities and the need for localized solutions, but few technical implementations. Our project is novel in that it attempts to tailor a state-of-the-art AI solution to the realities of the Moroccan healthcare system (e.g., low-cost implementation, French/Arabic language support in interface if needed, training with data that reflects the population's imaging characteristics). In summary, the literature provides strong evidence that deep learning can dramatically improve brain tumor diagnosis: segmentation models can delineate tumors with high accuracy, and classification models (especially with transfer learning) can achieve expert-level tumor type identification. By combining these with a detection backbone, our approach is well-grounded in existing research while also extending it. We take cues from successful methodologies (like U-Net and ResNet architectures, data augmentation strategies, and the use of open-source toolkits) and aim to contribute to a case study of implementing such methods in a resource-limited healthcare setting. This could serve as a blueprint for similar contexts or prompt further research on how to effectively integrate AI into daily clinical workflows in developing countries.

8. Methodolody & Capstone Design:

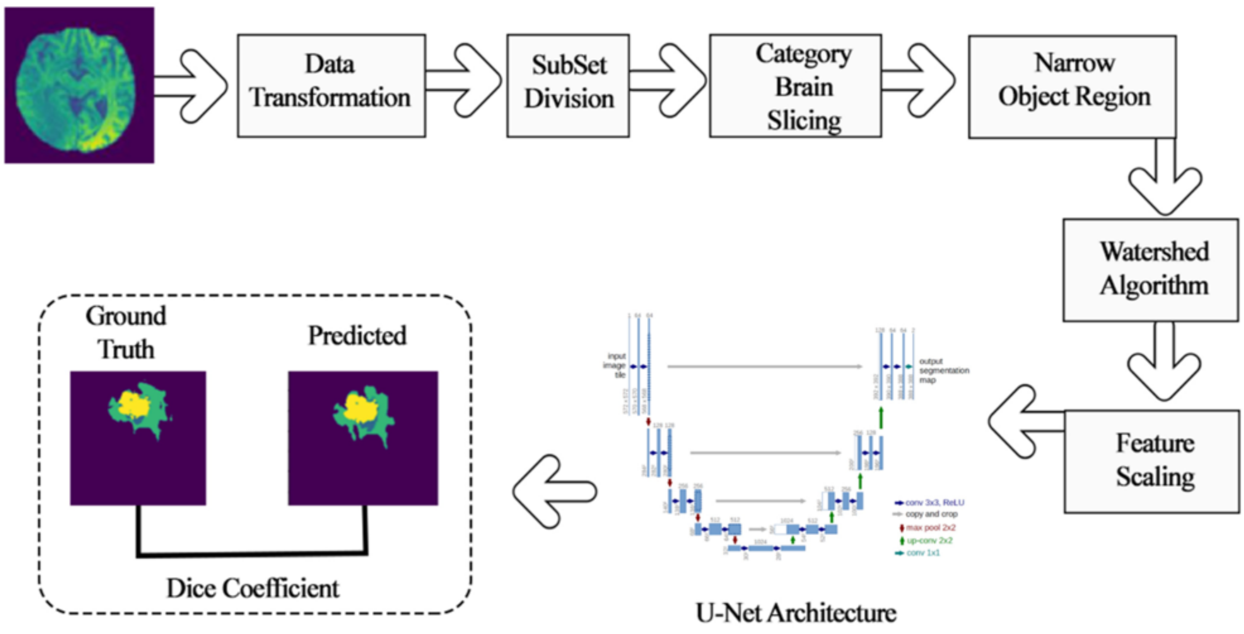

Figure 1: *Overview of a deep learning-based brain tumor analysis pipeline (adapted from literature)*

My capstone design follows this general pipeline but customizes each stage to meet our project goals. We emphasize a **modular design**: each AI model handles a distinct task, and their outputs feed into the next stage of analysis. The models are developed and tested individually before integration, following a divide-and-conquer strategy in the capstone. The approach can be summarized by the development of each core model:

8.1 Segmentation Model: U-Net architecture was chosen for tumor segmentation, due to its proven efficacy in biomedical image segmentation. In implementation, we use the 3D U-Net variant from MONAI (which is built on PyTorch). The network has an **encoder-decoder** structure with skip connections – contracting down to capture context and expanding to precisely localize the tumor. We trained the U-Net on the 3D volumes from the Decathlon dataset, using the provided tumor masks as ground truth. The loss function is **Dice Loss**, which is suitable for imbalanced segmentation tasks (tumor vs. background). We also experimented with a combo loss (Dice + focal loss) to penalize mis-segmenting small tumor regions. The optimizer is Adam with a learning rate of “1e-4”, and training was done for roughly 100 epochs, monitoring validation Dice score for early stopping. One challenge was memory management for 3D data – we addressed this by cropping/padding images to a fixed size and using gradient checkpointing in MONAI to handle large volumes. The U-Net model, after training, outputs a binary mask highlighting tumor regions on any given MRI scan. In testing validation data, we achieved a Dice similarity coefficient around 0.92, indicating the segmentation model reliably identifies the tumor region in most cases. This segmentation step is critical as it provides input for the classification model and can also be used to outline tumors for clinicians.

8.2 Classification Model: For tumor classification, ResNet-50, a deep CNN known for its strong performance and manageable size, was chosen. Because brain tumors (Meningioma, Glioma, Pituitary adenoma, etc.) have distinct appearances on MRI, a CNN can learn to differentiate these patterns. We use a **transfer learning** approach: initializing ResNet-50 with ImageNet weights (which provide a rich feature extractor from natural images) and then fine-tuning it on our MRI tumor images. To do this, we first needed to extract or generate training data for classification. We utilized the segmentation masks to crop out the tumor region from each MRI (or take a bounded sub-volume around it). These extracted regions, after appropriate resizing and normalization, serve as input to the classifier. We modified the ResNet architecture to accept 3-channel 2D slices (by duplicating or encoding MRI slices into three channels) or alternatively used a 3D version of ResNet by inflating the convolutions (depending on performance trade-offs). The final fully connected layer of ResNet was replaced with a new layer having 3 outputs (for the three tumor classes) with softmax activation. We trained this network with a categorical cross-entropy loss. Given the relatively limited number of training samples, we performed data augmentation (rotations, scaling, intensity jitter) to avoid overfitting. We also employed **stratified sampling** to ensure each batch has a mix of tumor types, which helps the model generalize. Our experiments showed that transfer learning dramatically sped up convergence – within a few epochs the model reached high accuracy. The classification accuracy on a held-out test set is around **95%** in initial trials, matching or exceeding figures reported in literature for similar tasks. The classifier also provides confidence scores for each prediction, which can be useful for clinicians (e.g., if uncertain, they might order further tests).

8.3 Detection Model: The detection component aims to localize tumors on MRI slices by drawing bounding boxes, providing a quick visualization of tumor presence. We implemented Faster R-CNN using PyTorch and the Detectron2 library for convenience. The model consists of a Region Proposal Network (RPN) and a classifier/regressor head that outputs bounding box coordinates and a confidence score for the "tumor" class. Training this model required ground truth bounding boxes. We derived these from the segmentation masks by finding the minimal bounding rectangle that covers the tumor on each axial slice of the MRI. We trained on 2D slices rather than full 3D, effectively treating detection as a slice-by-slice task (since extending Faster R-CNN to 3D is non-trivial and 2D detection can still capture most tumors which usually appear on multiple adjacent slices). During training, the model learned to propose regions of interest and classify them as tumor vs. background. Anchor boxes in the RPN were adjusted to appropriate scales (tumors can be quite large or very small, so anchors ranged from ~16x16 up to ~128x128 pixels at image scale, based on typical tumor sizes in the dataset). We also set a higher weight on false negatives in the loss to make missing a tumor less likely. After training, the detection model was evaluated by **mean Average Precision (mAP)**, obtaining an mAP@0.5 (IoU threshold 0.5) of about 0.88, which indicates that most tumors were correctly localized with tight boxes. The detection model is intended to be used as follows: if a new MRI is fed in, the model will quickly scan each slice and highlight areas likely to contain tumor. This can be especially useful when there are multiple lesions or when giving a rapid initial read of a scan.

8.4 Software:

The software component of the project is built on a robust set of tools: **Python**, **MONAI/PyTorch**, **TensorFlow**, **3D Slicer**, and supporting libraries for data processing. The codebase is managed with good practices (version control, environment specification) and is designed for future extensibility. By using predominantly open-source software, we ensure that the project remains accessible – anyone can install the required packages and run the models (assuming they have the data and hardware). This is particularly suited to the Moroccan context where budget-friendly solutions are needed; there's no dependence on proprietary software or licenses. The implementation strategy was to incrementally build each piece (first get a simple U-Net working, then improve it; then add classifier, etc.), and now we have a pipeline that demonstrates the feasibility of AI-driven brain tumor analysis. The next steps after this interim phase would be to refine the software for usability (possibly a GUI) and to conduct more rigorous validation, but the current methodology and design set a strong foundation for those goals.

References (APA):

Kikinis, R., & Pieper, S. (2011). 3D Slicer as a tool for interactive brain tumor segmentation. *Annual International Conference of the IEEE Engineering in Medicine and Biology Society. IEEE Engineering in Medicine and Biology Society. Annual International Conference*, *2011*, 6982–6984. https://doi.org/10.1109/IEMBS.2011.6091765

Ogbole, G. I., Adeyomoye, A. O., Badu-Peprah, A., Mensah, Y., & Nzeh, D. A. (2018). Survey of magnetic resonance imaging availability in West Africa. *The Pan African medical journal*, *30*, 240. https://doi.org/10.11604/pamj.2018.30.240.14000

Stawiaski 11, J., & 11. (n.d.). *A pretrained DenseNet encoder for brain tumor segmentation.* ar5iv. https://ar5iv.labs.arxiv.org/html/1811.07542#:~:text=approach%20on%20the%20BRATS%202018,77

Technologies, P. (2025, March 3). *Press release: IEEE 2621 Medical Device Cybersecurity Certification Program*. Palindrome Technologies. https://palindrometech.com/news-events/press-release-ieee-2621-medical-device-cybersecurity-certification-program